

Dataset Replay Report

Purpose: isolate **policy vs. execution** for the "cube not grasped, arm too high" failure. We replay the *exact recorded joint targets* from a training episode (no policy in the loop) and compare what the robot measures against what the dataset recorded.

- If replay-measured motion **matches** the dataset recording (and the cube is grasped) → execution/sim2real is faithful → the failure is in the **policy**.
- If the replay itself **diverges** (e.g. ends too high) → the gap is in **execution** (sim2real dynamics) → then re-run replay with TAM on to see if it closes that specific deviation.

Setup

Dataset	~/dataset/compliant_test_flip_pnp_softer/blocks_success
Control	joint PD, soft kp=40, kd=13 (pd_mode=0, matches the dataset)
Command	action.joint_target (absolute rad), replayed at 25 fps
Clamp	max_relative_movement = 20° (targets jump up to ~21°/step; see notes)
Gripper	inverted (dataset 0/1 → live 0=closed/1=open)
TAM	off for the first pass (bare execution gap), then on
Script	examples/inference/franka_replay.py
Plot	examples/inference/plot_replay.py

Data notes (from loading episode 0)

- 370 frames (14.8 s). pd_mode is all-0 (soft) as expected.
- **The terminal frame has action.is_valid=False and a garbage target** (130° jump on J4/J6) — dropped by the loader.
- Recorded **targets jump up to ~21°/step** while the recorded **observed motion is ≤1.8°/step** — the compliant controller ran heavily error-saturated (target far ahead, arm lagging). The 20° clamp therefore barely engages (~1 frame).

How to reproduce

```
# TAM off (bare execution gap):
python examples/inference/franka_replay.py
# then plot dataset vs measured:
python examples/inference/plot_replay.py tam_logs/<ts>

# TAM on (does it close the gap?): set ReplayConfig.tam = True, re-run, re-plot.
```

Outputs per run in tam_logs/<ts>/ : replay_frames.csv (per-frame q_sent / q_meas / q_ds + gripper), replay_{tam,notam}_controller_1khz.csv (1 kHz controller/TAM detail), replay_meta.json (episode/config provenance), replay_compare.png (the plot).

Results — episode 0, TAM off (20260821_161209)

The dataset is **SIMULATION** (TAMP planner: `outputs/compliant_test/flip_then_pnp` , domain/problem PDDL, per-episode domain-randomization metadata). So `q_ds` is the sim-achieved motion and `q_meas` is the real robot running the same targets — **this replay measures the sim2real execution gap directly**, with TAM off.

369 frames, 14.7 s.


Joint space — replay tracks the sim joint trajectory to a few degrees

$\text{RMS}(q_{\text{meas}} - q_{\text{ds}})$ per joint [deg]:


j0	j1	j2	j3	j4	j5	j6
1.32	2.93	0.69	1.13	0.74	0.82	0.61

$\max |q_{\text{meas}} - q_{\text{ds}}|$ per joint [deg]: [2.0, 5.1, 1.4, 1.9, 2.0, 1.8, 1.5] (worst is J1, the shoulder).

Replay ep0 (TAM off): dataset vs measured joint motion



Replay execution error (measured - dataset) per joint




EE space — small joint errors amplify to grasp-relevant cm-level EE error

Forward kinematics (FR3, **flange** frame — matches the dataset EE convention and the robot's `tcp_offset=identity`) of q_{meas} vs q_{ds} :

axis	RMS [mm]	max · [mm]
x	8.7	16.0
y	15.3	31.0
z (height)	36.3	64.7

Mean EE-Z(replay - sim) = **-29.7 mm** (replay runs ~3 cm *lower* on average), but the vertical error swings up to **6.5 cm** peak.



Verdict

- **Execution is NOT grasp-faithful.** Even replaying the planner's exact targets, the real robot's EE deviates from the sim by up to **6.5 cm (36 mm RMS), dominated by the vertical axis** — comparable to a cube's size, so enough to miss a grasp. The gap is small in joint space (1–3°) but amplifies through the kinematics.
- This is precisely the **sim2real execution gap TAM is meant to close**, and it is present with **TAM off**. So the "arm too high / cube not grasped" symptom is at least partly an **execution** problem, not purely the policy. (Note the sign: this replay averages ~3 cm *lower* than sim; the live policy looked *higher* — the two aren't directly comparable since the policy commands its own targets, but both point to an uncorrected vertical execution error.)

Results — episode 0, TAM on (20260821_162326) vs off

Same episode, TAM enabled (stable config: `tam_residual_clip=[...,0.5,0.5,0.5]` , `rate_limit=True` , `tam_history_pre_ratelimit=True`). TAM ran **stably** through the whole replay: `active=0.97` , OOD max **2.3σ**, wrist `J7 dq ≤0.73` rad/s, latent age median 99 ms — no limit cycle.

TAM improves joint tracking and EE x/y...

Joint RMS(`q_meas - q_ds`) [deg]:

	j0	j1	j2	j3	j4	j5	j6
TAM off	1.32	2.93	0.69	1.13	0.74	0.82	0.61
TAM on	0.59	2.38	0.43	1.28	0.35	0.44	0.41

EE deviation from sim [mm] (flange frame):

axis	RMS off	RMS on	max off	max on
x	8.7	7.6	16.0	14.2
y	15.3	9.6	31.0	21.8
z	36.3	33.9	64.7	51.9
3D	40.4	36.0		

TAM improves the **y** error (~37%), modestly improves x and the 3D total, and improves joint tracking on 6 of 7 joints — evidence it is closing part of the sim2real dynamics gap, and doing so stably.


...but it does NOT fix the vertical (grasp-critical) axis — it overcorrects it

The Z RMS barely moves (36 → 33 mm), and the **mean EE-Z bias flips sign**:


	mean EE-Z (replay – sim)
TAM off	–29.7 mm (arm ~3 cm too <i>low</i>)
TAM on	+31.5 mm (arm ~3 cm too <i>high</i>)

So TAM adds ~+60 mm of vertical correction and **overshoots** — turning a 3 cm *undershoot* into a 3 cm *overshoot*. **This matches the observed live symptom ("arm too high, cube not grasped")**: with TAM on, the EE sits ~3 cm above where the sim demo put it, which is enough to miss the cube.

Replay ep0: joint motion, TAM on vs off vs dataset



Replay ep0: EE trajectory & error, TAM on vs off (flange)



Verdict

- **TAM works and is stable** — it closes the dynamics gap in joint space and in EE x/y, without reigniting the wrist instability.
- **The grasp failure is a persistent ~3 cm vertical bias, not a TAM stability problem.** Both directions show a ~30 mm systematic EE-Z offset; TAM flips its sign (over-corrects) rather than removing it. A constant few-cm vertical bias that survives in both cases smells like a **calibration/model offset**, not a dynamics gap TAM can learn away.

Next steps

1. **Chase the ~3 cm vertical bias directly** (independent of TAM): check the **TCP offset / gripper mounting**, the **payload/load mass** used for gravity compensation, the **FK model vs. real** (this Z error appears in the bare soft controller too), and **table-height** calibration. A constant offset here would explain both signs.
2. **Understand TAM's vertical overshoot** — it adds ~2× the needed Z correction. Consider whether the shoulder/elbow (J1/J3, which drive EE-Z) residuals are too aggressive; a per-joint scale on those could center the Z error near zero.
3. Confirm grasp timing given the **inverted gripper** before attributing remaining failures.

*Caveat: gripper was replayed **inverted** (dataset → live convention); confirm open/close timing before trusting grasp outcome. EE numbers are FK (FR3 + Robotiq 2F-85 TCP) of measured vs. sim joints — a wrong TCP offset would bias the EE-Z comparison itself (see step 1).*

Dataset 2 — aug13_sim2real_kp40_pnp (pnp-only, kp40, sim)

Same two experiments on a second **simulation** dataset (pnp-only task, soft kp40). Episode 0, 262 frames.
Baseline 20260821_163425 (TAM off, rerun after the first attempt drove into the table), TAM-on 20260821_163735 .

TAM-on ran **stably**: active=0.95 , OOD max **1.8 σ** , wrist J7 $\dot{q} \leq 0.58$ rad/s, latent age median 96 ms — no limit cycle (wrist clip config unchanged).

The vertical-overshoot finding reproduces

Joint RMS($q_{meas} - q_{ds}$) [deg] — TAM improves 6 of 7 joints:

	j0	j1	j2	j3	j4	j5	j6
TAM off	1.37	2.65	0.69	1.08	0.82	0.75	0.86
TAM on	0.55	2.53	0.51	1.24	0.42	0.34	0.62


EE deviation from sim [mm] (flange):

axis	RMS off	RMS on	max off	max on
x	8.5	11.0	17.0	19.2
y	17.8	9.1	29.7	16.4
z	31.7	36.1	52.0	50.6
3D	37.3	38.8		


Mean EE-Z bias flips sign again — same as dataset 1:

	mean EE-Z (replay – sim)	EE-Z bottoms out at
TAM off	–27.7 mm (too low)	135 mm — 10 mm below sim (145 mm) → into the table
TAM on	+31.9 mm (too high)	182 mm — 37 mm above sim → cannot reach the cube

aug13 replay ep0: joint motion, TAM on vs off vs sim



aug13 replay ep0: EE trajectory & error, TAM on vs off (flange)



Cross-dataset conclusion

The vertical execution bias and TAM's response are **consistent across both datasets**:

	compliant_test	aug13_kp40_pnp
mean EE-Z, TAM off	-29.7 mm (low)	-27.7 mm (low)
mean EE-Z, TAM on	+31.5 mm (high)	+31.9 mm (high)

- **TAM off** → **arm ~3 cm too low** (drives into the table on the pnp dataset).
- **TAM on** → **arm ~3 cm too high** (bottoms out 37 mm above sim → misses the cube). This is exactly the live "arm too high, cube not grasped" symptom.
- TAM reliably improves joint tracking and the **y** axis and stays stable, but it **over-corrects the vertical axis by ~2x**, converting a -3 cm undershoot into a +3 cm overshoot rather than centering it. On aug13 the net EE 3D RMS is even slightly worse (37.3 → 38.8 mm) because the z overshoot dominates and x worsens.

The blocker is a systematic ~3 cm vertical offset, seen with TAM off, that TAM flips instead of removes. This is very likely a **calibration/model** issue (TCP/flange, payload/gravity comp, FK-vs-real, or table-height), not something TAM can learn away — pursue the vertical-bias hunt (report step 1) before expecting TAM to fix the grasp. A per-joint down-scale on the shoulder/elbow (J1/J3) residuals that drive EE-Z could also stop TAM from overshooting.

Dataset 2 with added payload — block mass in gripper


Same aug13 ep0 replay, but with a physical **block placed in the gripper** to add mass. TAM off 20260821_164347 , TAM on 20260821_164438 . (Note: the replay does not call `setLoad` , so the controller's gravity compensation does **not** know about the block — an unmodeled payload.)

EE-Z bias, with vs without the added mass:


run	mean EE-Z (replay – sim)	EE-Z bottoms at	joint RMS J1 [deg]
no-mass, TAM off	-27.7 mm	135 mm	2.65
mass, TAM off	-28.3 mm	135 mm	2.70
no-mass, TAM on	+31.9 mm	182 mm	2.53
mass, TAM on	+29.7 mm	184 mm	2.39

TAM-on with mass stayed **stable** (OOD max 1.8σ , J7 `dq` ≤ 0.58 , active 0.95).

aug13 replay ep0 WITH MASS: joint, TAM on vs off vs sim



aug13 replay ep0 WITH MASS: EE trajectory & error (flange)



Finding — the vertical bias is payload-independent

Adding the block **did not change** the results: EE-Z bias, the bottom-out height, and per-joint tracking are all within noise of the no-payload runs (off -28.3 vs -27.7 mm; on +29.7 vs +31.9 mm).

This is an informative **negative result**. If the ~3 cm vertical offset were driven by an **unmodeled payload / gravity-compensation** error, adding a block (which the controller's gravity comp ignores) should have made the soft arm sag noticeably *lower* with TAM off. It did not. So:

- **The ~3 cm vertical bias is NOT a gravity/load-modeling problem** — it is robust to payload.
- That points the calibration hunt (report step 1) toward a **payload-independent kinematic/model offset**: TCP/flange definition, FK-model-vs-real, mounting, or table-height reference — rather than load mass or gravity compensation.
- TAM's vertical **overshoot to ~+3 cm reproduces with payload too**, unchanged.